Chapter 2 Message-Passing Computing

Basics of Message-Passing Programming
Programming Options

Programming a message-passing multicomputer can be achieved by

1. Designing a special parallel programming language

2. Extending the syntax/reserved words of an existing sequential high-level language to handle message passing

3. Using an existing sequential high-level language and providing a library of external procedures for message passing

We will concentrate upon the third option. Necessary to say explicitly what processes are to be executed, when to pass messages between concurrent processes, and what to pass in the messages.

Two primary methods are needed in this form of a message-passing system:

1. A method of creating separate processes for execution on different computers

2. A method of sending and receiving messages

Single Program Multiple Data (SPMD) model

Different processes are merged into one program. Within the program are control statements that will customize the code; i.e. select different parts for each process.

Figure 2.1   Single program, multiple data operation.

Multiple Program Multiple Data (MPMD) Model

Separate programs written for each processors. Master-slave approachusually taken whereby a single processor executes a master process and other processes are started from within the master process.
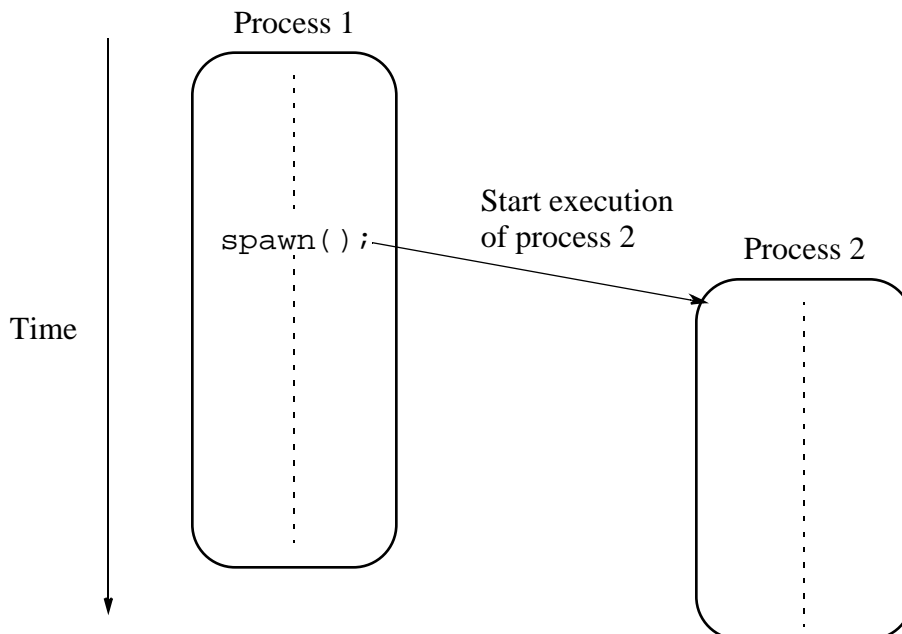
Figure 2.2   Spawning a process.

Basic Send and Receive Routines

Process 1
Process 2

```
        x                           Y
send(&x, 2);      Movement
                   of data
                              recv(&y, 1);
```

Figure 2.3   Passing a message between processes using `send()` and `recv()` library calls.
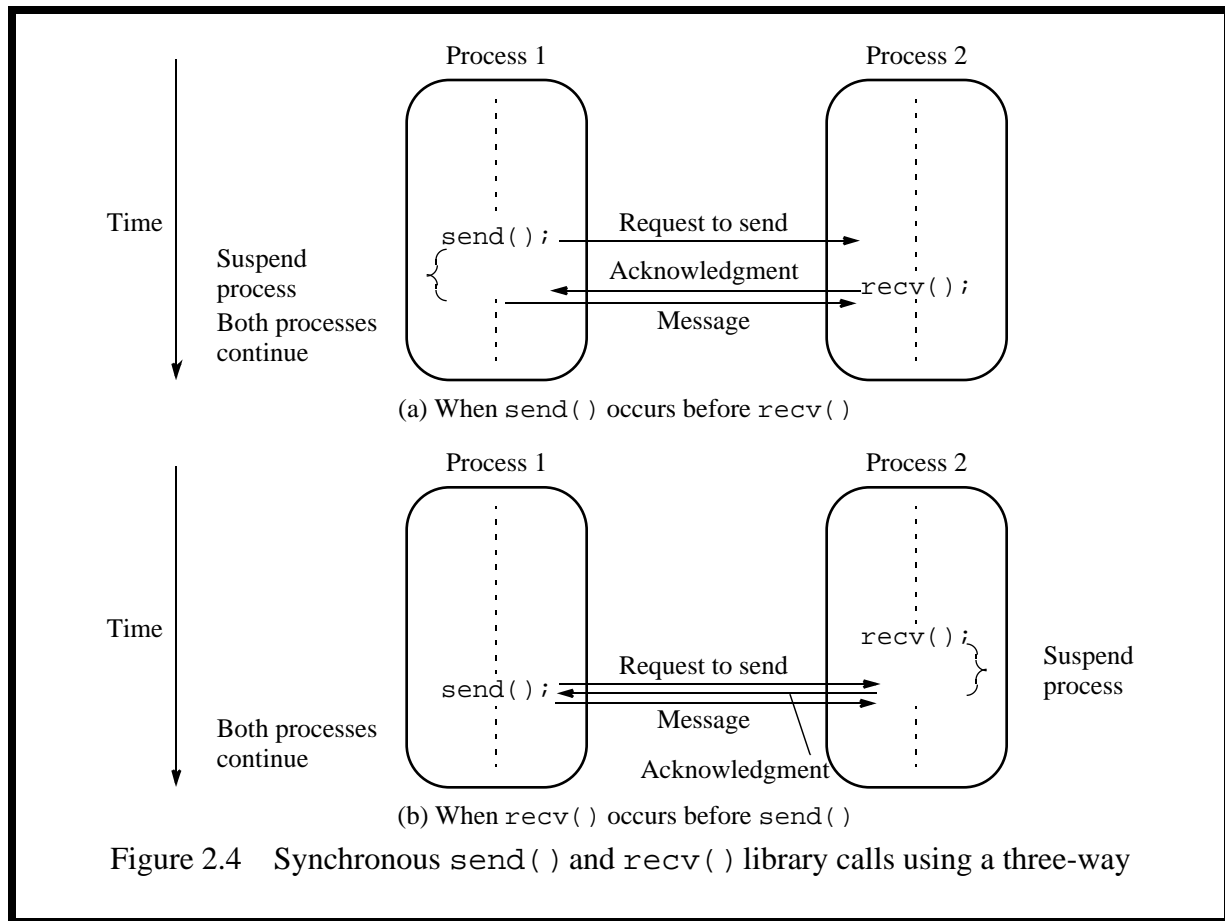
Synchronous Message Passing

Routines that actually return when the message transfer has been completed.

Do not need message buffer storage. A synchronous send routine could wait until the complete message can be accepted by the receiving process before sending the message.

A synchronous receive routine will wait until the message it is expecting arrives.

Synchronous routines intrinsically perform two actions: They transfer data and they synchronize processes.

Suggest some form of signaling, such as a three-way protocol:

(a) When send() occurs before recv()

(b) When recv() occurs before send()

Figure 2.4   Synchronous send() and recv() library calls using a three-way

Blocking and Nonblocking Message Passing

Blocking - has been used to describe routines that do not return until the transfer is completed.

The routines are "blocked" from continuing.

In that sense, the terms *synchronous* and *blocking* were synonymous.

Non-blocking - has been used to describe routines that return whether or not the message had been received.

The terms *blocking* and *nonblocking* redefined in systems such as MPI:

MPI Definitions of Blocking and Non-Blocking

Blocking - return after their local actions complete, though the message transfer may not have been completed.

Non-blocking - return immediately. Assumed that the data storage being used for the transfer is not modified by the subsequent statements prior to the data storage being used for the transfer, and it is left to the programmer to ensure this.

How message-passing routines can return before the message transfer has been completed

Generally, a *message buffer* needed between source and destination to hold message:
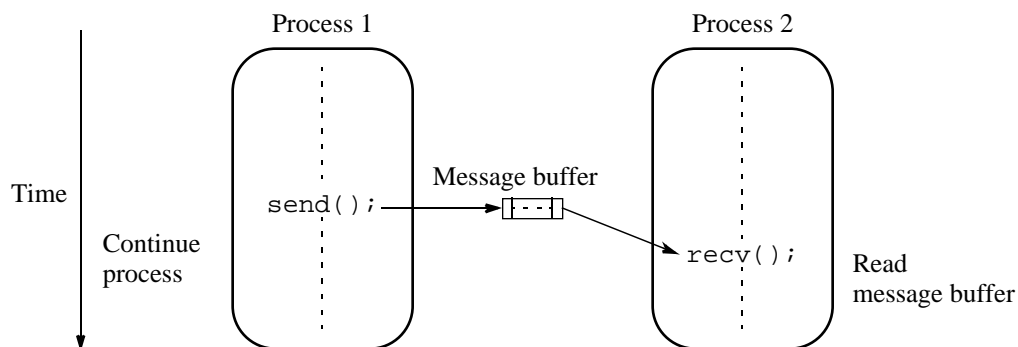

Figure 2.5   Using a message buffer.

For send routine, once the local actions have been completed and the message is safely on its way, the process can continue with subsequent work. Buffers can only be of finite length and a point could be reached when the send routine is held up because all the available buffer space has been exhausted.

Message Tag

Used to differentiate between different types of messages being sent.

Example

To send a message, `x,` with message tag 5 from a source process, 1, to a destination process, 2, and assign to `y`, we might have

```
send(&x, 2, 5);
```

in the source process and

```
recv(&y, 1, 5);
```

in the destination process. The message tag is carried within the message.

If special type matching is not required, a *wild card* message tag is used, so that the `recv()` will match with any `send()`.

Broadcast

Sending the same message to all the processes concerned with the problem.

M*ulticast* - sending the same message to a defined group of processes.

Figure 2.6   Broadcast operation.

Scatter

Sending each element of an array of data in the root to a separate process. The contents of the ith location of the array is sent to the ith process.


Figure 2.7   Scatter operation.

Gather

Having one process collect individual values from a set of processes.


Figure 2.8   Gather operation.

Reduce

Gather operation combined with a specified arithmetic or logical operation. Example, the values could be gathered and then added together by the root:
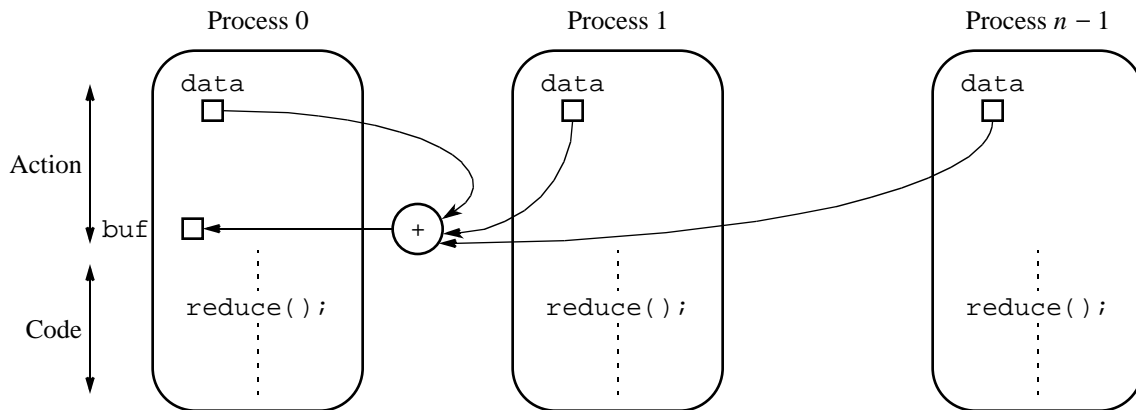
Process 0   Process 1   Process $n-1$

data   data   data

Action

buf

Code

reduce();   reduce();   reduce();

Figure 2.9   Reduce operation (addition).

Using Workstation Clusters - Software Tools

PVM (Parallel Virtual Machine) - Perhaps the first widely adopted attempt at using a workstation cluster as a multicomputer platform developed by Oak Ridge National Laboratories.

Provides for a software environment for message passing between homogeneous or heterogeneous computers and has a collection of library routines that the user can employ with C or FORTRAN programs. Available at no charge.

MPI (Message Passing Interface) - standard developed by group of academics and industrial partnersto foster more widespread use and portability. Several free implementations exist

PVM

The programmer decomposes the problem into separate programs. Each program is written in C (or Fortran) and compiled to run on specific types of computers in the network.

The set of computers used on a problem first must be defined prior to running the programs.

The most convenient way of doing this is by creating a list of the names of the computers available in a *hostfile*. The hostfile is then read by PVM.

The routing of messages between computers is done by PVM daemon processes installed by PVM on the computers that form the virtual machine:


Figure 2.10   Message passing between workstations using PVM.

Figure 2.11   Multiple processes allocated to each processor (workstation).

Basic Message-Passing Routines

All PVM send routines are nonblocking (or asynchronous in PVM terminology) while PVM receive routines can be either blocking (synchronous) or nonblocking. Uses a *message tag* (`msgtag`). Both message tag and source wild cards available.

`pvm_psend()`and `pvm_precv()`

If data being sent is a list of items of the same data type, the PVM routines pvm_psend() and pvm_precv() can be used.

Parameter in pvm_psend() points to an array of data in the source process to be sent, and a parameter in pvm_precv() points to where to store the received data.

Process 1   Send buffer   Process 2

Array holding data

Array to receive data

Pack

`pvm_psend();`

Continue process

`pvm_precv();` Wait for message

Figure 2.12   `pvm_psend()` and `pvm_precv()` system calls.

Sending Data Composed of Various Types

Data packed into a send buffer prior to sending data. Receiving process must unpack its receive buffer according to format in which it was packed. Specific packing and unpacking routines for each datatype.

Process_1

Process_2

```
pvm_initsend();

pvm_pkint( … &x …);
pvm_pkstr( … &s …); -
pvm_pkfloat( … &y …);
pvm_send(process_2 … );
```

Send buffer

Message

x
s
Y

```
pvm_recv(process_1 …);
pvm_upkint( … &x …);
pvm_upkstr( … &s …);
pvm_upkfloat(… &y … );
```

Receive buffer

Figure 2.13   PVM packing messages, sending, and unpacking.

Broadcast, Multicast, Scatter, Gather, and Reduce

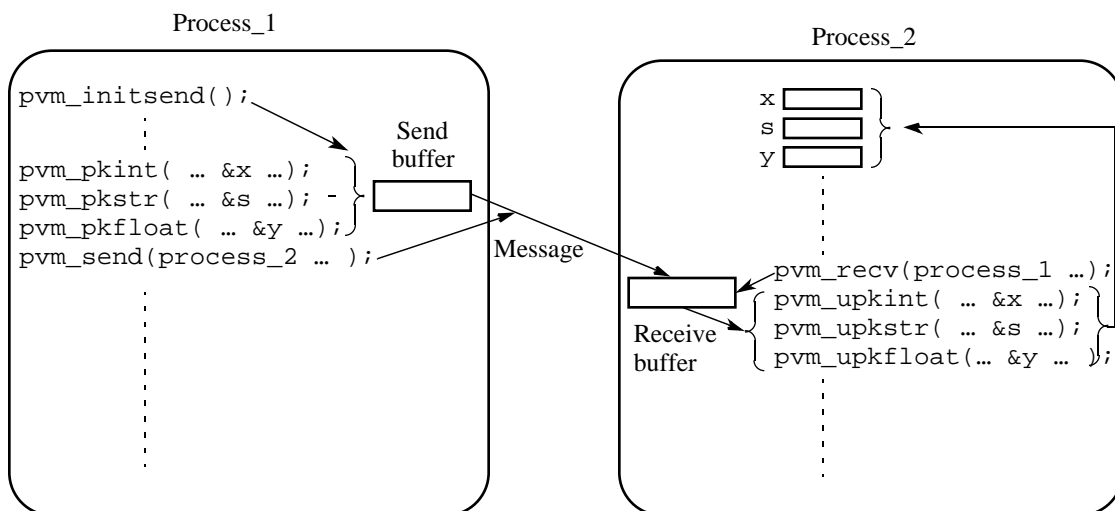Broadcast, scatter, gather, and reduce operations (`pvm_bcast()`, `pvm_scatter()`, `pvm_gather()`, and `pvm_reduce()`) used with group of processes.

A process joins named group by calling `pvm_joingroup()`.

The `pvm_bcast()`, when called, would send a message to each member of the named group.

Similarly, `pvm_gather()` would collect values from each member of the named group.

The PVM multicast operation, `pvm_mcast()`, is not a group operation.

```
#include <stdio.h>              Master
#include <stdlib.h>
#include <pvm3.h>
#define SLAVE "spsum"
#define PROC 10
#define NELEM 1000
main() {
   int mytid,tids[PROC];
   int n = NELEM, nproc = PROC;
   int no, i, who, msgtype;
   int data[NELEM],result[PROC],tot=0;                        Slave
   char fn[255];
   FILE *fp;                                    #include <stdio.h>
   mytid=pvm_mytid();/*Enroll in PVM */         #include "pvm3.h"
                                                #define PROC 10
/* Start Slave Tasks */                         #define NELEM 1000
  no=
   pvm_spawn(SLAVE,(char**)0,0,"",nproc,tids);  main() {
  if (no < nproc) {                                int mytid;
    printf("Trouble spawning slaves \n");          int tids[PROC];
    for (i=0; i<no; i++) pvm_kill(tids[i]);         int n, me, i, msgtype;
    pvm_exit(); exit(1);                            int x, nproc, master;
  }                                                int data[NELEM], sum;

/* Open Input File and Initialize Data */          mytid = pvm_mytid();
  strcpy(fn,getenv("HOME"));
  strcat(fn,"/pvm3/src/rand_data.txt");        /* Receive data from master */
  if ((fp = fopen(fn,"r")) == NULL) {            msgtype = 0;
    printf("Can't open input file %s\n",fn);      pvm_recv(-1, msgtype);
    exit(1);                                      pvm_upkint(&nproc, 1, 1);
  }                                              pvm_upkint(tids, nproc, 1);
  for(i=0;i<n;i++)fscanf(fp,"%d",&data[i]);      pvm_upkint(&n, 1, 1);
                                                 pvm_upkint(data, n, 1);
/* Broadcast data To slaves*/
  pvm_initsend(PvmDataDefault);                /* Determine my tid */
  msgtype = 0;
```

```
/* Open Input File and Initialize Data */              mytid = pvm_mytid();
  strcpy(fn,getenv("HOME"));
  strcat(fn,"/pvm3/src/rand_data.txt");             /* Receive data from master */
  if ((fp = fopen(fn,"r")) == NULL) {                 msgtype = 0;
    printf("Can't open input file %s\n",fn);          pvm_recv(-1, msgtype);
    exit(1);                                           pvm_upkint(&nproc, 1, 1);
  }                                                    pvm_upkint(tids, nproc, 1);
  for(i=0;i<n;i++)fscanf(fp,"%d",&data[i]);           pvm_upkint(&n, 1, 1);
                                                       pvm_upkint(data, n, 1);
/* Broadcast data To slaves*/
  pvm_initsend(PvmDataDefault);                     /* Determine my tid */
  msgtype = 0;                                         for (i=0; i<nproc; i++)
  pvm_pkint(&nproc, 1, 1);                               if(mytid==tids[i])
  pvm_pkint(tids, nproc, 1);                               {me = i;break;}
  pvm_pkint(&n, 1, 1);
  pvm_pkint(data, n, 1);            Broadcast data  /* Add my portion Of data */
  pvm_mcast(tids, nproc, msgtag);                     x = n/nproc;
                                                       low = me * x;
                                                       high = low + x;
/* Get results from Slaves*/                           for(i = low; i < high; i++)
  msgtype = 5;                                            sum += data[i];
  for (i=0; i<nproc; i++){
    pvm_recv(-1, msgtype);                          /* Send result to master */
    pvm_upkint(&who, 1, 1);                            pvm_initsend(PvmDataDefault)
    pvm_upkint(&result[who], 1, 1);                    pvm_pkint(&me, 1, 1);
    printf("%d from %d\n",result[who],who);  Receive results  pvm_pkint(&sum, 1, 1);
  }                                                    msgtype = 5;
                                                       master = pvm_parent();
/* Compute global sum */                               pvm_send(master, msgtype);
  for (i=0; i<nproc; i++) tot += result[i];
  printf ("The total is %d.\n\n", tot);            /* Exit PVM */
                                                       pvm_exit();
  pvm_exit(); /* Program finished. Exit PVM */         return(0);
  return(0);                                        }
}
```

Figure 2.15   Sample PVM program.

MPI
Process Creation and Execution

Purposely not defined and will depend upon the implementation.

Only static process creation is supported in MPI version 1. All the processes must be defined prior to execution and started together. Use SPMD model of computation.

Communicators

Defines the *scope* of a communication operation.

Processes have ranks associated with the communicator.

Initially, all processes enrolled in a "universe" called MPI_COMM_WORLD, and each process is given a unique rank, a number from 0 to $n - 1$, where there are n processes.

Other communicators can be established for groups of processes.

Using the SPMD Computational Model

```
main (int argc, char *argv[])
{
  MPI_Init(&argc, &argv);
   .
   .
  MPI_Comm_rank(MPI_COMM_WORLD, &myrank);/* find process rank */
  if (myrank == 0)
    master();
  else
    slave();
   .
   .
  MPI_Finalize();
}
```

where `master()` and `slave()` are procedures to be executed by the master process and slave process, respectively.

Global and Local Variables

Any global declarations of variables will be duplicated in each process.

Variables that are not to be duplicated will need to be declared within code only executed by that process.

Example

```
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);/* find process rank */
if (myrank == 0) {       /* process 0 actions/local variables */
  int x, y;
     .
     .
} else if (myrank == 1) {/* process 1 actions/local variables */
  int x, y;
     .
     .
}
```

Here, x and y in process 0 are different local variables from x and y in process 1.
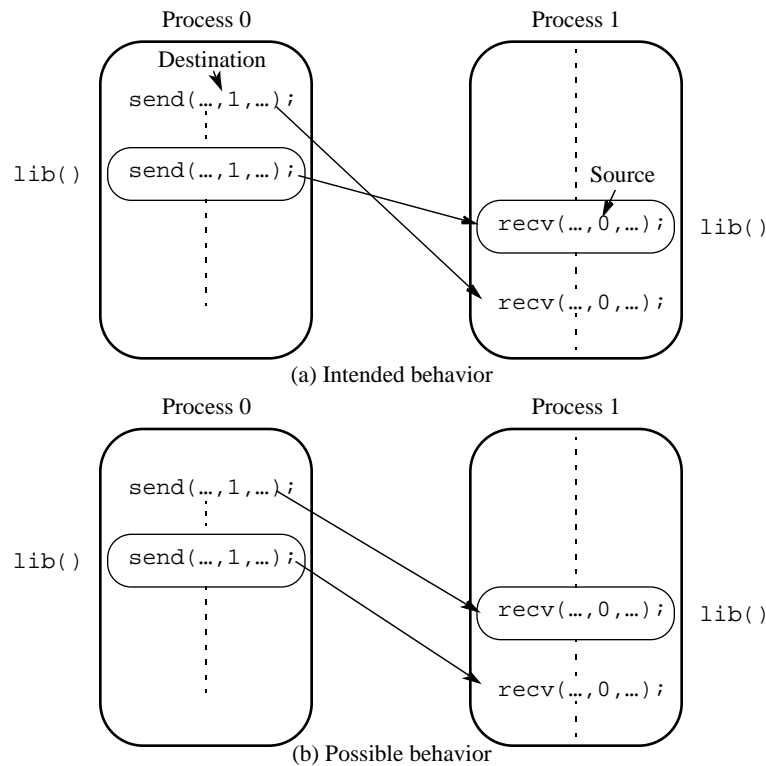
Figure 2.16   Unsafe message passing with libraries.

MPI Solution

Communicators - used in MPI for all point-to-point and collective MPI message-passing communications.

A communicator is a *communication domain* that defines a set of processes that are allowed to communicate between themselves.

In this way, the communication domain of the library can be separated from that of a user program.

Each process has a rank within the communicator, an integer from 0 to $n - 1$, where there are n processes.

Communicator Types

Intracommunicator - for communicating within a group

Intercommunicator - for communication between groups.

A process has a unique *rank* in a group (an integer from 0 to *m* − 1, where there are *m* processes in the group). A process could be a member of more than one group.

Default intracommunicator - `MPI_COMM_WORLD`, exists as the first communicator for all the processes existing in the application.New communicators are created based upon existing communicators. A set of MPI routines exists for forming communicators.

Point-to-Point Communication

Message tags are present, and wild cards can be used in place of the tag (`MPI_ANY_TAG`) and in place of the source in receive routines (`MPI_ANY_SOURCE`).

PVM style packing and unpacking data is generally avoided by the use of an MPI datatype being defined in the send/receive parameters together with the source or destination of the message.

Blocking Routines

Return when they are locally complete - when the location used to hold the message can be used again or altered without affecting the message being sent.

A blocking send will send the message and return. This does not mean that the message has been received, just that the process is free to move on without adversely affecting the message.

The general format of parameters of the blocking send is

MPI_Send(buf, count, datatype, dest, tag, comm)

Address of
send buffer

Datatype of
each item

Message tag

Number of items
to send

Rank of destination
process

Communicator

The general format of parameters of the blocking receive is

MPI_Recv(buf, count, datatype, src, tag, comm, status)

Address of
receive buffer

Datatype of
each item

Message tag

Status after operation

Maximum number
of items to receive

Rank of source
process

Communicator

Example

To send an integer *x* from process 0 to process 1,

```
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);/* find process rank */
if (myrank == 0) {
  int x;
  MPI_Send(&x, 1, MPI_INT, 1, msgtag, MPI_COMM_WORLD);
} else if (myrank == 1) {
  int x;
  MPI_Recv(&x, 1, MPI_INT, 0, msgtag, MPI_COMM_WORLD, status);
}
```

Nonblocking Routines

Nonblocking send - `MPI_Isend()`, will return "immediately" even before source location is safe to be altered.

Nonblocking receive - `MPI_Irecv()`, will return even if there is no message to accept.

Formats

```
MPI_Isend(buf, count, datatype, dest, tag, comm, request)
MPI_Irecv(buf, count, datatype, source, tag, comm, request)
```

Completion detected by `MPI_Wait()` and `MPI_Test()`.

`MPI_Wait()` waits until the operation has actually completed and will return then.
`MPI_Test()` returns with a flag set indicating whether operation completed at that time.

These routines need to know whether the particular operation has completed, which is determined by accessing the `reques`t parameter.

Example

To send an integer x from process 0 to process 1 and allow process 0 to continue,

```
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);/* find process rank */
if (myrank == 0) {
  int x;
  MPI_Isend(&x, 1, MPI_INT, 1, msgtag, MPI_COMM_WORLD, req1);
  compute();
  MPI_Wait(req1, status);
} else if (myrank == 1) {
  int x;
  MPI_Recv(&x, 0, MPI_INT, 1, msgtag, MPI_COMM_WORLD, status);
}
```

Send Communication Modes
Standard Mode Send

Not assumed that corresponding receive routine has started. Amount of buffering not defined by MPI. If buffering is provided, send could complete before receive reached.

Buffered Mode

Send may start and return before a matching receive. Necessary to specify buffer space via routine `MPI_Buffer_attach()` - removed with `MPI_Buffer_detach()`.

Synchronous Mode

Send and receive can start before each other but can only complete together.

Ready Mode

Send can only start if matching receive already reached, otherwise error. Use with care.

Each of the four modes can be applied to both blocking and nonblocking send routines. Only the standard mode is available for the blocking and nonblocking receive routines. Any type of send routine can be used with any type of receive routine.

Collective Communication

Involves set of processes, defined by an intra-communicator. Message tags not present.

Broadcast and Scatter Routines

The principal collective operations operating upon data are

```
MPI_Bcast()          - Broadcast from root to all other processes
MPI_Gather()         - Gather values for group of processes
MPI_Scatter()        - Scatters buffer in parts to group of processes
MPI_Alltoall()       - Sends data from all processes to all processes
MPI_Reduce()         - Combine values on all processes to single value
MPI_Reduce_scatter() - Combine values and scatter results
MPI_Scan()           - Compute prefix reductions of data on processes
```

Example

To gather items from the group of processes into process 0, using dynamically allocated memory in the root process, we might use

```
int data[10];                   /*data to be gathered from processes*/
        .
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);   /* find rank */
if (myrank == 0) {
    MPI_Comm_size(MPI_COMM_WORLD, &grp_size); /*find group size*/
    buf = (int *)malloc(grp_size*10*sizeof(int));/*allocate memory*/
}
MPI_Gather(data,10,MPI_INT,buf,grp_size*10,MPI_INT,0,MPI_COMM_WORLD);
```

Note that `MPI_Gather()` gathers from all processes, including the root.

Barrier

As in all message-passing systems, MPI provides a means of synchronizing processes by stopping each one until they all have reached a specific "barrier" call.

```
#include "mpi.h"
#include <stdio.h>
#include <math.h>
#define MAXSIZE 1000
void main(int argc, char *argv)
{
   int myid, numprocs;
   int data[MAXSIZE], i, x, low, high, myresult, result;
   char fn[255];
   char *fp;
   MPI_Init(&argc,&argv);
   MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
   MPI_Comm_rank(MPI_COMM_WORLD,&myid);
   if (myid == 0) {              /* Open input file and initialize data */
     strcpy(fn,getenv("HOME"));
     strcat(fn,"/MPI/rand_data.txt");
     if ((fp = fopen(fn,"r")) == NULL) {
        printf("Can't open the input file: %s\n\n", fn);
        exit(1);
     }
     for(i = 0; i < MAXSIZE; i++) fscanf(fp,"%d", &data[i]);
   }
   /* broadcast data */
   MPI_Bcast(data, MAXSIZE, MPI_INT, 0, MPI_COMM_WORLD);
 /* Add my portion Of data */
   x = n/nproc;
   low = myid * x;
   high = low + x;
   for(i = low; i < high; i++)
     myresult += data[i];
   printf("I got %d from %d\n", myresult, myid);
 /* Compute global sum */
   MPI_Reduce(&myresult, &result, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
   if (myid == 0) printf("The sum is %d.\n", result);
   MPI_Finalize();
}
```
Figure 2.17   Sample MPI program.

Pseudocode Constructs

To send message consisting of an integer x and a float y, from the process called master
to the process called slave, assigning to a and b, we simply write in the master process

    send(&x, &y, P_slave);

and in the slave process

    recv(&a, &b, P_master);

where x and a are integers and y and b are floats. x will be copied to a, and y copied to b.

ith process given notation P_i, and a tag may be present i.e.,

    send(&x, P_2, data_tag);

sends **x** to process 2, with the message tag **data_tag**.

Locally blocking send() and recv() written as given. Other forms with prefixes; i.e.,

    ssend(&data1, P_destination);   /* Synchronous send */

Parallel Execution Time

Two parts: a computation part, say t_{comp}, and a communication part, say t_{comm}; i.e.,

$$t_p = t_{comp} + t_{comm}$$

Computation time estimated in a similar way to that of a sequential algorithm.

Communication Time

As a first approximation, we will use

$$t_{comm} = t_{startup} + n t_{data}$$

where $t_{startup}$ is the *startup time*, sometimes called the *message latency* - essentially time to send a message with no data. Startup time is assumed constant.

The term t_{data} is the transmission time to send one data word, also assumed a constant, and there are *n* data words.

Figure 2.18   Theoretical communication time.

Important Note on Interpretation of Equations

Only intended to give a starting point to how an algorithm might perform in practice.

Parallel execution time, t_p, normalized in units of an arithmetic operation, which will depend upon computer system.

We might find that the computation requires m computational steps so that

$$t_{comp} = m$$

Since we are measuring time in units of computational steps, the communication time has to be measured in the same way.

All data types are assumed to require the same time.

Suppose q messages are sent, each containing n data items. We have

$$t_{comm} = q(t_{startup} + nt_{data})$$

Latency Hiding

A way to ameliorate situation of significant message communication times is to overlap communication with subsequent computations.

Nonblocking send routines provided particularly to enable latency hiding.

Latency hiding can also be achieved by mapping multiple processes on a processor and use a time-sharing facility that switches for one process to another when the first process is stalled because of incomplete message passing or otherwise.

Relies upon an efficient method of switching from one process to another. *Threads* offer an efficient mechanism.

Time Complexity

As with sequential computations, a parallel algorithm can be evaluated through the use of time complexity (notably the Ο notation — "order of magnitude," big-oh).

Start with an estimate of the number of the computational steps, considering all arithmetic and logical operations to be equal and ignoring other aspects of the computation such as computational tests.

An expression of the number of computational steps is derived, often in terms of the number of data items being handled by the algorithm.

The Ο notation

$f(x) = Ο(g(x))$ if and only if there exists positive constants, c and x_0, such that $0 \le f(x) \le cg(x)$ for all $x \ge x_0$

where $f(x)$ and $g(x)$ are functions of x.

For example, if $f(x) = 4x^2 + 2x + 12$, the constant $c = 6$ would work with the formal definition to establish that $f(x) = O(x^2)$, since $0 < 4x^2 + 2x + 12 \le 6x^2$ for $x \ge 3$.

Alternative functions for $g(x)$ that will satisfy definition. Use the function that grows least for $g(x)$.

notation - upper bound

$f(x) = O(g(x))$ if and only if there exists positive constants c_1, c_2, and x_0 such that $0 \le c_1 g(x) \le f(x) \le c_2 g(x)$ for all $x \ge x_0$.

If $f(x) = \Theta(g(x))$, it is clear that $f(x) = O(g(x))$ is also true.

notation - lower bound

$f(x) = \Omega(g(x))$ if and only if there exists positive constants c and x_0 such that $0 \le cg(x) \le f(x)$ for all $x \ge x_0$.

It follows from this definition that $f(x) = 4x^2 + 2x + 12 = \Omega(x^2)$

We can read $O()$ as "grows at most as fast as" and $\Omega()$ as "grows at least as fast as."

Example

The execution time of a sorting algorithm often depends upon the original order of the numbers to be sorted. It may be that it requires at least $n \log n$ steps, but could require n^2 steps for n numbers depending upon the order of the numbers. This would be indicated by a time complexity of $\Omega(n \log n)$ and $O(n^2)$.

The figure at top:

Figure 2.19   Growth of function $f(x) = 4x^2 + 2x + 12$.

Time Complexity of a Parallel Algorithm

Time complexity analysis hides lower terms - t_{comm} has time complexity of (n).

Time complexity of t_p will be sum of complexity of computation and communication.

Example

To add n numbers on two computers, where each computer adds $n/2$ numbers together, and the numbers are initially all held by first computer. Second computer submits its result to first computer for adding the two partial sums together. Several phases:

   1.   Computer 1 sends $n/2$ numbers to computer 2.
   2.   Both computers add $n/2$ numbers simultaneously.
   3.   Computer 2 sends its partial result back to computer 1.
   4.   Computer 1 adds the partial sums to produce the final result.

Computation (for steps 2 and 4):

$$t_{comp} = n/2 + 1$$

Communication (for steps 1 and 3):

$$t_{comm} = (t_{startup} + n/2 t_{data}) + (t_{startup} + t_{data}) = 2t_{startup} + (n/2 + 1)t_{data}$$

The computational complexity is (n). The communication complexity is (n). The overall time complexity is (n).

Cost-Optimal Algorithms

A *cost-optimal* (or *work-efficient* or *processor-time optimality*) algorithm is one in which the cost to solve a problem is proportional to the execution time on a single processor system (using the fastest known sequential algorithm); i.e.,

$$\text{Cost} = t_p \times n = k \times t_s$$

where *k* is a constant.

Given time complexity analysis, we can say that a parallel algorithm is cost-optimal algorithm if

(Parallel time complexity) × (number of processors) = sequential time complexity

Example

Suppose the best known sequential algorithm for a problem has time complexity of (*n* log *n*). A parallel algorithm for the same problem that uses *n* processes and has a time complexity of (log *n*) is cost optimal, whereas a parallel algorithm that uses n^2 processors and has time complexity of (1) is not cost optimal.

Time Complexity of Broadcast/Gather
Broadcast on a Hypercube Network

Consider a three-dimensional hypercube. To broadcast from node 000 to every other node, 001, 010, 011, 100, 101, 110 and 111, an efficient algorithm is

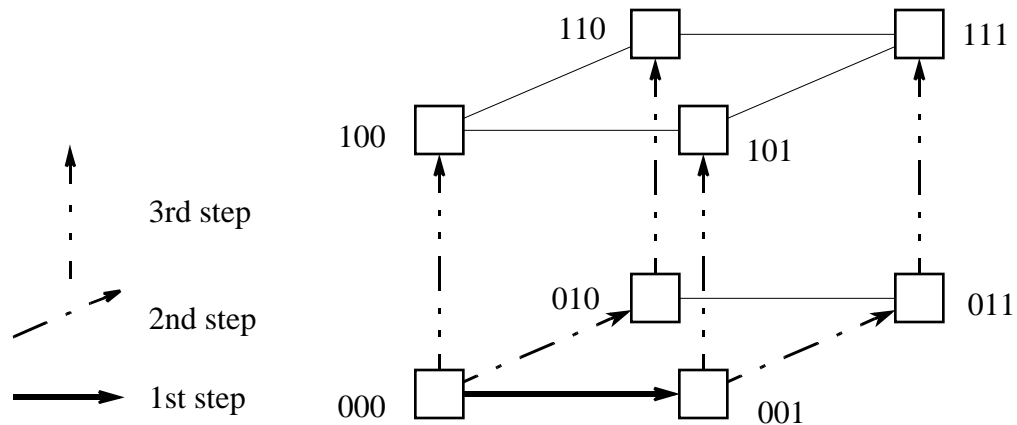|	Node	Node
1st step:	000	001
2nd step:	000	010
	001	011
3rd step:	000	100
	001	101
	010	110
	011	111

Figure 2.20   Broadcast in a three-dimensional hypercube.

The time complexity for a hypercube system will be   (log *n*), using this algorithm, which is optimal because the *diameter* of a hypercube network is log *n*. It is necessary at least to use this number of links in the broadcast to reach the furthest node.


Figure 2.21   Broadcast as a tree construction.

Gather on a Hypercube Network

The reverse algorithm can be used to gather data from all nodes to, say, node 000; i.e., for a three-dimensional hypercube,

	Node	Node
1st step:	100	000
	101	001
	110	010
	111	011
2nd step:	010	000
	011	001
3rd step:	001	000

In the case of gather, the messages become longer as the data is gathered, and hence the time complexity is increased over (log n).

Broadcast on a Mesh Network

Send message across top row and down each column as it reaches top of that column.

Steps

Figure 2.22   Broadcast in a mesh.

Optimal in terms of number of steps because same as diameter of mesh.

Broadcast on a Workstation Cluster

Broadcast on a single Ethernet connection can be done using a single message that is read by all the destinations on the network simultaneously


Figure 2.23   Broadcast on an Ethernet network.


Figure 2.24   1-to-*N* fan-out broadcast.

Source

Figure 2.25   1-to-*N* fan-out broadcast on a tree structure.

Debugging and Evaluating Parallel Programs
Visualization Tools

Programs can be watched as they are executed in a *space-time diagram* (or *process-time diagram*):
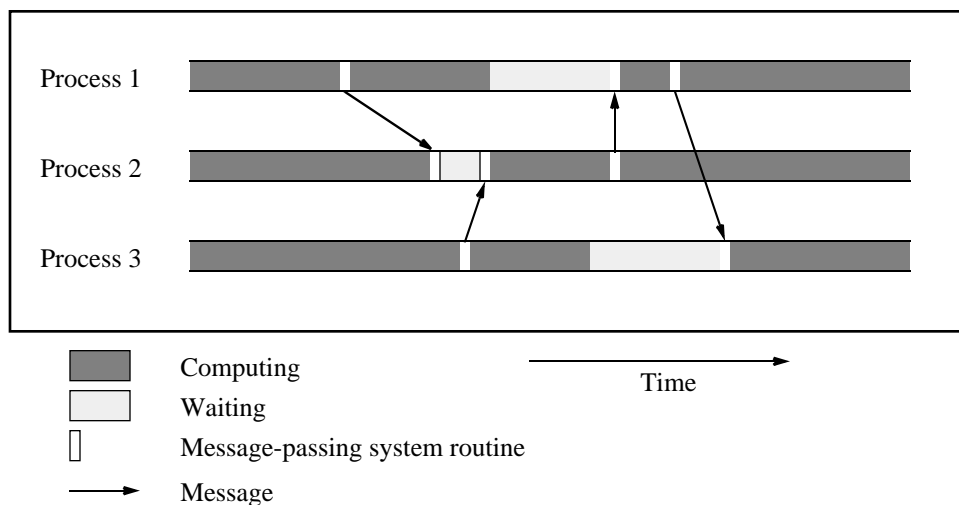

Figure 2.26   Space-time diagram of a parallel program.

PVM has a visualization tool called XPVM.

Implementations of visualization tools are available for MPI. An example is the Upshot program visualization system.

Debugging Strategies

Geist et al. (1994a) suggest a three-step approach to debugging message-passing programs:

1. If possible, run the program as a single process and debug as a normal sequential program.

2. Execute the program using two to four multitasked processes on a single computer. Now examine actions such as checking that messages are indeed being sent to the correct places. It is very common to make mistakes with message tags and have messages sent to the wrong places.

3. Execute the program using the same two to four processes but now across several computers. This step helps find problems that are caused by network delays related to synchronization and timing.

Evaluating Programs Empirically
Measuring Execution Time

To measure the execution time between point L1 and point L2 in the code, we might have

a construction such as

```
        .
    L1: time(&t1);                      /* start timer */
        .
        .
    L2: time(&t2);                      /* stop timer */
        .
        elapsed_time = difftime(t2, t1); /* elapsed_time = t2 - t1 */
        printf("Elapsed time = %5.2f seconds", elapsed_time);
```

MPI provides the routine **MPI_Wtime()** for returning time (in seconds).

Communication Time by the Ping-Pong Method

One process, say P_0, is made to send a message to another process, say P_1.

Immediately upon receiving the message, P_1 sends the message back to P_0.

```
    P0
      .
L1: time(&t1);
    send(&x, P1);
    recv(&x, P1);
L2: time(&t2);
    elapsed_time = 0.5 * difftime(t2, t1);
    printf("Elapsed time = %5.2f seconds", elapsed_time);
      .

    P1
      .
    recv(&x, P0);
    send(&x, P0);
      .
```
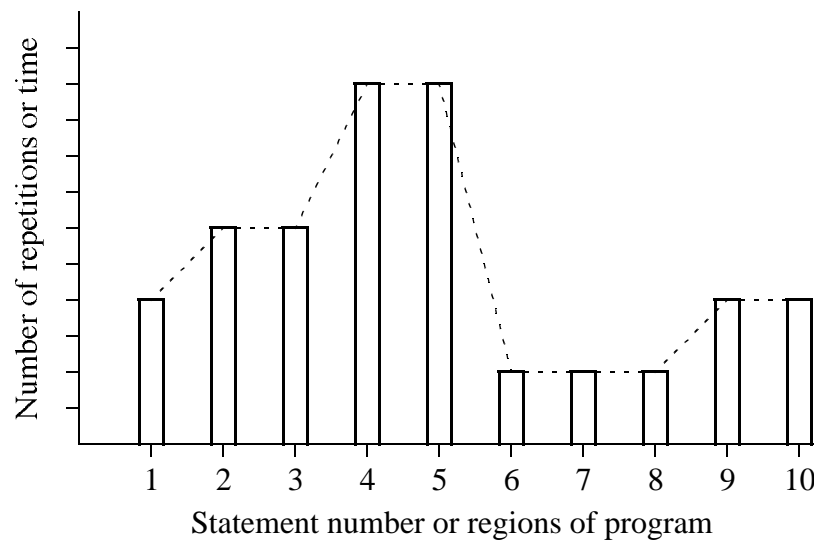
Profiling

A *profile* is a histogram or graph showing time spent on different parts of program:

Figure 2.27   Program profile.